

Sample test construction:

In testing the performance and reliability of each model, I used a sample test that divided up the work into 20 threads (chosen arbitrarily), performed 10000000 swaps (I found that in testing values even larger than 10000000, *GetNSet* would often fail to terminate. I believe this to be related to a change in the array values that would prevent a true return value in swap, causing the program to fail to terminate), a max range of 100, and 10 random values for the array. I created these testing large testing parameters to allow for greater variations in my swaps and a more pronounced multithreading result (as more threads access the same elements), resulting in more distinct performance/reliability results across each model. The performance measure for each model is measured as an average over all successful runs out of 5 total runs; the reliability measure is measured as an average total sums from all 5 runs (out of original value 450).

Performance/reliability measurements:

MODEL	TIME(NS/TRANSITION)	SUM MISMATCH
Null	194.4406 (0 fails)	N/A
Synchronized	3995.71 (0 fails)	N/A
Unsynchronized	859.064 (2 fails)	867.74333
GetNSet	3129.525 (3 fails)	922
BetterSafe	1726.08 (0 fails)	N/A
BetterSorry	273.089 (0 fails)	N/A

*SynchronizedState*: *SynchronizedState* is DRF as the *Synchronized* keyword prevents thread interference and memory consistency errors from multiple threads reading and writing to the same location in memory.

*UnsynchronizedState*: *UnsynchronizedState* is not DRF as illustrated by the sum mismatch errors that were produced by the sample test parameters above; this is because, without the keyword *Synchronized*, this allows for multiple threads to access the elements of *UnsynchronizedState*, such as in the *swap* method which reads and modifies the elements *value[i]* and *value[j]*. The implementation for *UnsynchronizedState* was relatively easy.

*GetNSet*: *GetNSet* is not DRF as illustrated by the sum mismatch errors that were produced by the sample test parameters above; this is in due part to the *get* methods within the *swap* function. Although these *get* functions are themselves atomic, their sequential orderings/execution is not threadsafe. In implementing *GetNSet*, I first had trouble with implementing the *current* method. After finding out

that I could not simply return the *AtomicIntegerArray* *value* by simple type-casting it as (byte), I experimented with several methods before type casting each element first and then putting them into an instantiated return list. Then, I used the *set* functions native to the *AtomicIntegerArray* library to increment and decrement the *value* variables.

*BetterSafe*: *BetterSafe* is DRF as illustrated by: (1) the non-existence of sum mismatch errors and (2) the use of lock mechanisms in my code. In implementing a *private final* lock variable that is locked and unlocking in *swap*, the function that made *GetNSet* not DRF, each *value* element is accessed/modified by one thread at a time. This preserves 100% reliability in running the code and performs better than *Synchronized*, although in documentation, both prevent multiple concurrent accesses to the same area in memory by multiple threads. I had initially thought that this lock implementation would make *BetterSafe* slower than *SynchronizedState* due to the additional overhead; however, with the dramatic performance improvements, I think that it is because *Reentrantlock* objects do not have the block locking structure that *Synchronized* does, but rather, an unstructured locking approach that is more appropriate for the increment/decrement/read/write operations of the program.

*BetterSorry*: In my implementation of *BetterSorry*, I was not able to provided an implementation that was faster than *BetterSafe* and, at the same time, not DRF. At first, I had tried an implementation that took advantage of the atomic nature of *AtomicIntegerArray* (inspired by *GetNSet*, using atomic increment/decrement functions) based on the intuition that atomic objects had less overhead than both locks and synchronization (from documentation). However, this approach ended up being less fast than *BetterSafe*. My implementation does, however, perform better than *BetterSafe* and *GetNSet* while maintaining a high level of reliability (possibly 100%). In *BetterSorry*, I created a *private ReentrantLock* array, capitalizing upon its *isHeldByCurrentThread* function in my *swap* function to prevent concurrent race conditions; values are only modified if they are held by the same thread. Although I could not find any conclusive tests to prove that my implementation is not DRF, one race condition that could possibly exist would be concurrent accesses to the *value* objects in the *swap* method (as explained above).

Given my data, I would use *BetterSorry* for GDI applications (due to its high performance/reliability ratio). However, given that there may be a data race, I would use *BetterSafe*. However, one must note that these values are subjective to the running system; the results are relative to each other and not concrete as I got different values when running my program at different times.